

Using AI to grow a private practice

THE WORKSHOP

A full day for private practice owners, run as four 90-minute sessions that build on each other, using the free tiers of Claude, ChatGPT, Gemini, or Copilot. Participants worked on their own practice throughout and left with content, a communication plan, and a draft business plan.

Foundations

Learning how the model works, what the context window holds, what to keep out of it, and how to structure a prompt.

Creating content

Building a persona for one client group and writing for that person in your own voice.

Getting found and trusted

Reviewing your website as a client would, then writing an AI-use policy, staff guide, and communication plan.

Growing the business

Drafting a three-year plan with Deep Research or an interview, then stress-testing it.

CONCEPTS

Context window

Everything the model can see in one conversation: your messages, its replies, and anything you've pasted or uploaded. The model works from this and from what it learned in training, and every new conversation starts empty.

[READ MORE →](#)

Sycophancy

The model's tendency to agree with you: to praise the plan you brought and confirm the idea you favour. Agreement from the model isn't evidence that you're right. It will also drop a correct answer when you push back.

[READ MORE →](#)

Persona

A written description of one client group: who they are, what brings them in, and what stops them from moving forward. You supply what you know, ask the model to fill the gaps, and correct what it invents. Content written for one person reads better than content written for everyone.

KEY TAKEAWAYS

01

Give context before asking for output

"Write me a blog post on back pain" produces something that could be about any practice. The same request, preceded by a paragraph about your practice and your clients, produces something you could use. Write that paragraph once and keep it. When you're not sure what the model needs to know, ask it to interview you.

02

Keep the clinical judgement yours

The model drafts; you decide. Never paste in anything you don't have permission to share, including client records, staff information, and other people's writing. Describe the situation without the identifying detail, and ask the model to reason through it with you.

03

Give every piece of content somewhere to go

A newsletter, a website change, a staff guide for the conversation that starts "ChatGPT told me...", and an AI-use policy your clients can read are all drafts until each has a channel, an owner, and a way of knowing whether it worked. That's the communication plan.

04

Ask the model to argue against its own plan

A three-year plan written by a system that knows nothing about your town, your lease, or your staff will be confidently wrong in places. The most useful prompt for a plan like this is also the simplest: *you wrote this plan; now argue against it*. Use it on anything you're about to act on.

FURTHER READING

Context engineering for educators: a prompt is the instruction, and context is everything else the model works with. It was written for universities, but the point holds for a practice: what you give the model matters more than your phrasing.

Building AI personas for professional practice: why "you are an expert physiotherapist" gets you the average of everything ever written, and how to describe your own approach so the model works within it.

The hidden inefficiency in how we work with AI: why it pays to match the model to the task, since some handle reasoning well, some long documents, and some quick routine work.

Organising your notes for AI: how to structure what you know so a model can use it without losing what makes it yours, and why plain text you can paste into any tool is worth keeping.