SIDRA MEDICINE | PEDIATRIC NEUROSCIENCE CONFERENCE | ONLINE | 2026-10-30 | MICHAEL ROWE

Working with AI agents in health systems

THE TALK

An invited 40-minute talk for nurses and allied health professionals working in paediatric neuroscience, given online. It asks what AI agents can now do, and how well the whole system of care works for the child now that clinicians, families and AI all take part in it.

What AI agents can do now

AI tools now remember, plan, act and check their work, and the evidence has moved from test cases to real patients. How people ask still decides what they get back.

How we're responding

The usual responses leave clinicians watching change happen. Asking what clinicians and AI can reach together keeps the child at the centre.

Care as a system

Clinical reasoning has always been shared across people, records and protocols. The family's AI, and yours, are now part of that system.

Families' own agents

Families arrive having already asked. Helping them ask well is a new kind of health literacy, and we'll work it out with them.

What this asks of you

Setting up what the AI does, noticing when the system drifts, and answering for decisions you didn't wholly produce.

CONCEPTS

Sanctuary strategies

Four common responses to AI in practice: denial, retreat to what only humans can do, restriction, and resignation. Each assumes a stable boundary between human and machine work. The boundary keeps moving, so each strategy gives up ground with every new model release.

Distributed cognition

Clinical reasoning is spread across people and tools: the team, the records, the protocols, the family, and now AI. Your judgement shapes how well that whole system thinks, including what the AI sees and what gets checked.

READ MORE →

AI agent

An AI system that keeps context across time, plans a sequence of steps, acts and checks its own work. Families are building up their own record of their child's history and worries in conversations with AI that you may never see.

READ MORE →

KEY TAKEAWAYS

01

Ask the next family you see how they're using AI

Make it an open, routine question: what have you asked it about your child, and what did it tell you? You'll hear what the family already believes, and they'll know it's something you want to hear about.

02

Help families ask their AI better questions

Suggest open questions in place of yes-or-no ones, the context worth giving (age, diagnosis, medicines, what's changed), what to leave out (names, dates of birth), and that they bring back what the AI said.

FURTHER READING

We're comparing AI chatbot health advice to the wrong thing: why warning patients off AI applies a standard we don't apply anywhere else in the system, and what a better response looks like.

Reliability of LLMs as medical assistants for the general public: Bean et al. (2026) in *Nature Medicine*. The AI alone named the right condition 95% of the time; people using it managed under 35%.

Context sovereignty for AI-supported learning: why what an AI knows about you shapes what it can do for you, and who should control that context.

Making sense of AI in clinical practice: an earlier talk for physiotherapists that goes further into the context clinicians bring to AI.

Handout produced for an invited talk in the Neuroscience Nursing and Allied Health Symposium at the Sidra Medicine Pediatric Neuroscience Conference. Slides, references, and notes.

CC BY 4.0